

```
In [85]: import pandas as pd
import random
import seaborn as sns
sns.set(color_codes=True)
from sklearn.preprocessing import LabelEncoder
import numpy as np
from pandas_datareader import data
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import r2_score
from sklearn.datasets import load_boston
from sklearn.model_selection import GridSearchCV
```

```
In [2]: pip install pandas-datareader
```

```
Requirement already satisfied: pandas-datareader in c:\users\priya\anaconda3\lib\site-packages (0.10.0)
Requirement already satisfied: requests>=2.19.0 in c:\users\priya\anaconda3\lib\site-packages (from pandas-datareader) (2.25.1)
Requirement already satisfied: pandas>=0.23 in c:\users\priya\anaconda3\lib\site-packages (from pandas-datareader) (1.2.4)
Requirement already satisfied: lxml in c:\users\priya\anaconda3\lib\site-packages (from pandas-datareader) (4.6.3)
Requirement already satisfied: numpy>=1.16.5 in c:\users\priya\anaconda3\lib\site-packages (from pandas>=0.23->pandas-datareader) (1.19.5)
Requirement already satisfied: python-dateutil>=2.7.3 in c:\users\priya\anaconda3\lib\site-packages (from pandas>=0.23->pandas-datareader) (2.8.1)
Requirement already satisfied: pytz>=2017.3 in c:\users\priya\anaconda3\lib\site-packages (from pandas>=0.23->pandas-datareader) (2021.1)
Requirement already satisfied: six>=1.5 in c:\users\priya\anaconda3\lib\site-packages (from python-dateutil>=2.7.3->pandas>=0.23->pandas-datareader) (1.15.0)
Requirement already satisfied: chardet<5,>=3.0.2 in c:\users\priya\anaconda3\lib\site-packages (from requests>=2.19.0->pandas-datareader) (4.0.0)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in c:\users\priya\anaconda3\lib\site-packages (from requests>=2.19.0->pandas-datareader) (1.26.4)
Requirement already satisfied: idna<3,>=2.5 in c:\users\priya\anaconda3\lib\site-packages (from requests>=2.19.0->pandas-datareader) (2.10)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\priya\anaconda3\lib\site-packages (from requests>=2.19.0->pandas-datareader) (2020.12.5)
Note: you may need to restart the kernel to use updated packages.
```

```
In [3]: df = pd.read_csv('heart-statlog_csv.csv', na_values = ['?'])
```

In [4]: df.head()

Out[4]:

	age	sex	chest	resting_blood_pressure	serum_cholestorol	fasting_blood_sugar	resting_electro
0	70	1	4		130	322	0
1	67	0	3		115	564	0
2	57	1	2		124	261	0
3	64	1	4		128	263	0
4	74	0	2		120	269	0

In [5]: df.shape

Out[5]: (270, 14)

In [6]: df.isna().sum()

Out[6]:

age	0
sex	0
chest	0
resting_blood_pressure	0
serum_cholestorol	0
fasting_blood_sugar	0
resting_electrocardiographic_results	0
maximum_heart_rate Achieved	0
exercise_induced_angina	0
oldpeak	0
slope	0
number_of_major_vessels	0
thal	0
class	0
dtype: int64	

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 270 entries, 0 to 269
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   age              270 non-null    int64  
 1   sex              270 non-null    int64  
 2   chest             270 non-null    int64  
 3   resting_blood_pressure  270 non-null  int64  
 4   serum_cholestorol  270 non-null    int64  
 5   fasting_blood_sugar 270 non-null    int64  
 6   resting_electrocardiographic_results 270 non-null  int64  
 7   maximum_heart_rate_achieved 270 non-null    int64  
 8   exercise_induced_angina 270 non-null    int64  
 9   oldpeak            270 non-null    float64 
 10  slope              270 non-null    int64  
 11  number_of_major_vessels 270 non-null    int64  
 12  thal               270 non-null    int64  
 13  class              270 non-null    object  
dtypes: float64(1), int64(12), object(1)
memory usage: 29.7+ KB
```

In [8]: le=LabelEncoder()  
label=le.fit\_transform(df["class"])

In [9]: le.classes\_

Out[9]: array(['absent', 'present'], dtype=object)

In [10]: Data=df.drop("class",axis='columns')

In [11]: Data.head()

Out[11]:

	age	sex	chest	resting_blood_pressure	serum_cholestorol	fasting_blood_sugar	resting_electro
0	70	1	4		130	322	0
1	67	0	3		115	564	0
2	57	1	2		124	261	0
3	64	1	4		128	263	0
4	74	0	2		120	269	0

In [12]: Data["class"]=label

In [13]: Data

Out[13]:

g_blood_sugar	resting_electrocardiographic_results	maximum_heart_rate_achieved	exercise_induced_angina
0	2	109	
0	2	160	
0	0	141	
0	0	105	
0	2	121	
...	...	...	...
1	0	162	
0	0	173	
0	2	153	
0	0	148	
0	2	108	

In [14]: `import matplotlib.pyplot as plt  
plt.style.use('fivethirtyeight')`

In [15]: `df['class'].unique()`

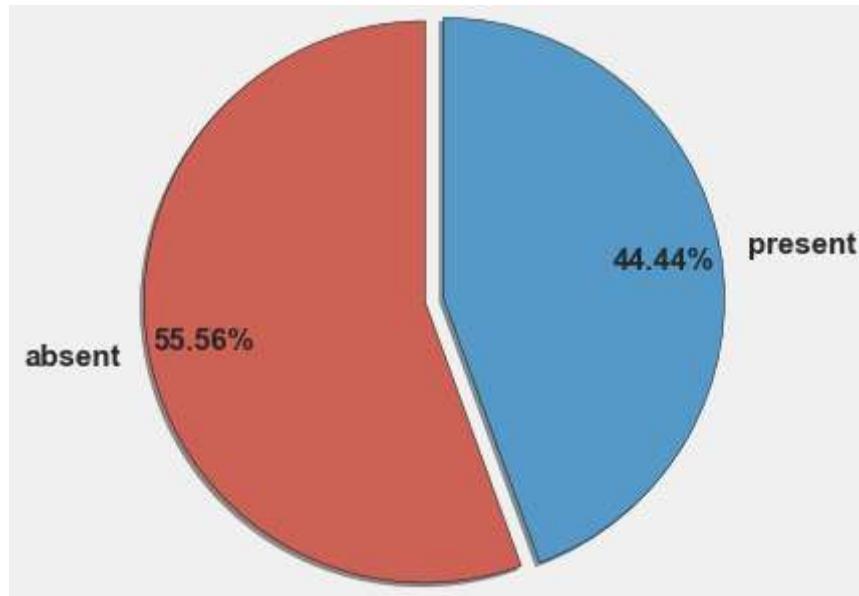
Out[15]: `array(['present', 'absent'], dtype=object)`

In [16]: `work_count = df['class'].value_counts().tolist()`

In [17]: `work_label = df['class'].value_counts().index`

In [18]: `colors = ['#CD6155', '#5499C7', '#AF7AC5', '#48C9B0', '#52BE80', '#F4D03F']`

```
In [19]: plt.pie(work_count, labels = work_label,
               autopct = '%1.2f%%',
               colors = colors[:3],
               wedgeprops = {'edgecolor':'k'},
               textprops = {'fontweight':'bold', 'size':15},
               shadow = True,
               explode = [0.1, 0],
               startangle = 90,
               pctdistance = 0.8,
               radius=1.5)
plt.show()
plt.savefig(r'E:\overleaf charts\barchart.png')
```

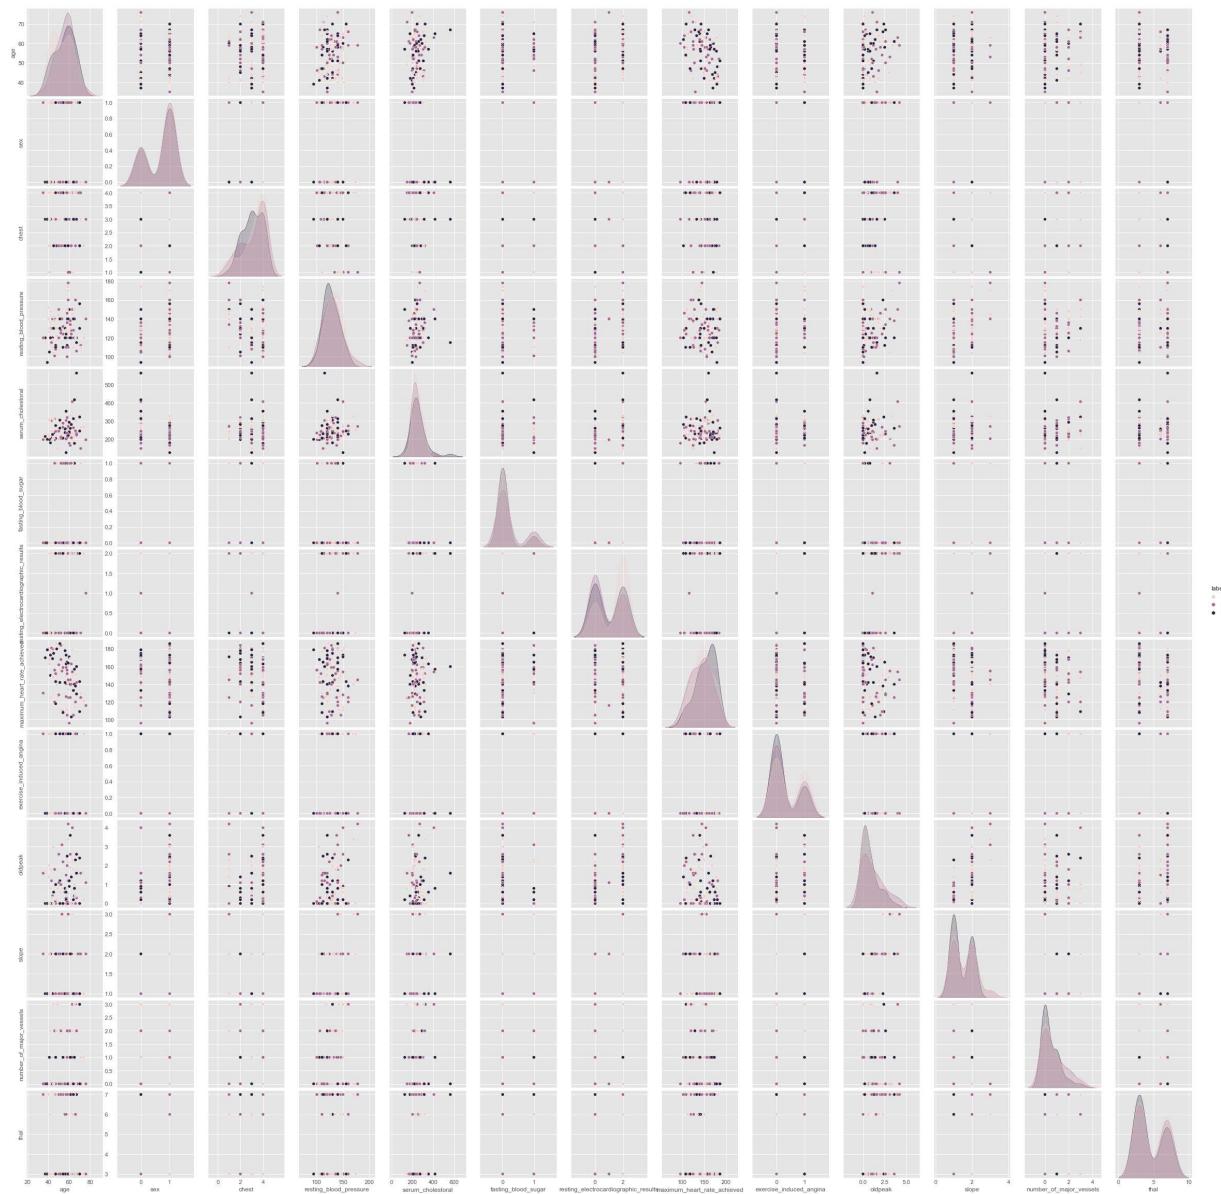


<Figure size 432x288 with 0 Axes>

## PSO

```
In [68]: from sklearn.datasets import make_classification
X, y = make_classification(n_samples=100, n_features=14, n_classes=3,
                           n_informative=7, n_redundant=1, n_repeated=2,
                           random_state=1)
```

```
In [69]: df['labels'] = pd.Series(y)  
sns.pairplot(df, hue='labels');
```



In [22]: pip install pyswarms

```
Requirement already satisfied: pyswarms in c:\users\priya\anaconda3\lib\site-packages (1.3.0)
Requirement already satisfied: tqdm in c:\users\priya\anaconda3\lib\site-packages (from pyswarms) (4.59.0)
Requirement already satisfied: numpy in c:\users\priya\anaconda3\lib\site-packages (from pyswarms) (1.19.5)
Requirement already satisfied: future in c:\users\priya\anaconda3\lib\site-packages (from pyswarms) (0.18.2)
Requirement already satisfied: attrs in c:\users\priya\anaconda3\lib\site-packages (from pyswarms) (20.3.0)
Requirement already satisfied: matplotlib>=1.3.1 in c:\users\priya\anaconda3\lib\site-packages (from pyswarms) (3.3.4)
Requirement already satisfied: pyyaml in c:\users\priya\anaconda3\lib\site-packages (from pyswarms) (5.4.1)
Requirement already satisfied: scipy in c:\users\priya\anaconda3\lib\site-packages (from pyswarms) (1.6.2)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.3 in c:\users\priya\anaconda3\lib\site-packages (from matplotlib>=1.3.1->pyswarms) (2.4.7)
Requirement already satisfied: pillow>=6.2.0 in c:\users\priya\anaconda3\lib\site-packages (from matplotlib>=1.3.1->pyswarms) (8.2.0)
Requirement already satisfied: cycler>=0.10 in c:\users\priya\anaconda3\lib\site-packages (from matplotlib>=1.3.1->pyswarms) (0.10.0)
Requirement already satisfied: python-dateutil>=2.1 in c:\users\priya\anaconda3\lib\site-packages (from matplotlib>=1.3.1->pyswarms) (2.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\priya\anaconda3\lib\site-packages (from matplotlib>=1.3.1->pyswarms) (1.3.1)
Requirement already satisfied: six in c:\users\priya\anaconda3\lib\site-packages (from cycler>=0.10->matplotlib>=1.3.1->pyswarms) (1.15.0)
Note: you may need to restart the kernel to use updated packages.
```

In [23]: RANDOM\_SEED =1

```
In [60]: random.seed(RANDOM_SEED)
np.random.seed(RANDOM_SEED)

SWARM_SIZE=30
NUM_ITER=1000

NUM_PATIENTS=40
MAX_AGE=max(df['age'])

age=df['age']
```

```
In [25]: def f_per_particle(m, max_age):
    """Computes for the objective function per particle

    Inputs
    -----
    m : numpy.ndarray
        Binary mask that can be obtained from BinaryPSO, will
        be used to mask features.
    alpha: float (default is 0.5)
        Constant weight for trading-off classifier performance
        and number of features

    Returns
    -----
    numpy.ndarray
        Computed objective function
    """

    total_features = 14
    # Get the subset of the features from the binary mask
    if np.count_nonzero(m) == 0:
        X_subset = X
    else:
        X_subset = X[:,m==1]
    # Perform classification and store performance in P
    classifier.fit(X_subset, y)
    P = (classifier.predict(X_subset) == y).mean()
    # Compute for the objective function
    j = (alpha * (1.0 - P)
         + (1.0 - alpha) * (1 - (X_subset.shape[1] / total_features)))

    return j

def f_per_particle(m, alpha):
    """Computes for the objective function per particle

    Inputs
    -----
    m : numpy.ndarray
        Binary mask that can be obtained from BinaryPSO, will
        be used to mask features.
    alpha: float (default is 0.5)
        Constant weight for trading-off classifier performance
        and number of features

    Returns
    -----
    numpy.ndarray
        Computed objective function
    """

    total_features = 13
    # Get the subset of the features from the binary mask
    if np.count_nonzero(m) == 0:
        X_subset = X
    else:
        X_subset = X[:,m==1]
    # Perform classification and store performance in P
```

```
classifier.fit(X_subset, y)
P = (classifier.predict(X_subset) == y).mean()
# Compute for the objective function
j = (alpha * (1.0 - P)
     + (1.0 - alpha) * (1 - (X_subset.shape[1] / total_features)))

return j
```

```
In [26]: import numpy as np
import seaborn as sns
import pandas as pd

# Import PySwarms
import pyswarms as ps
```

```
In [27]: from tensorflow.keras.optimizers import SGD
```

```
In [28]: def f(x,max_age=MAX_AGE):
    n_particles=x.shape[0]
    j=[f_per_particle(x[i],max_age) for i in range(n_particles)]
    return np.array(j)
```

```
In [29]: def run_model(options):
    optimizer = ps.discrete.BinaryPSO(n_particles=SWARM_SIZE, dimensions=15, opti
    cost, pos = optimizer.optimize(f, iters=1000)
    return cost,pos,optimizer.cost_history
```

```
In [30]: def plot_history(history):
    plt.style.use('ggplot')
    plt.rcParams['ytick.right']=True
    plt.rcParams['ytick.labelright']=True
    plt.rcParams['ytick.left']=False
    plt.rcParams['ytick.labelleft']=False
    plt.rcParams['font.family']='Arial'

    plt.ylim([min[history]-1000,max[history]+500])
    plt.title('Cost History')
    plt.plot(history)
```

```
In [31]: checkpoint_state=np.random.get_state()
```

## INITIAL MODEL

```
In [32]: import keras
from keras.models import Sequential
from keras.layers import Dense
from keras.models import load_model

#empty model
classifier = Sequential()
```

```
In [33]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
```

```
In [34]: inputs = keras.Input(shape=(784,), name="digits")
x = layers.Dense(64, activation="relu", name="dense_1")(inputs)
x = layers.Dense(64, activation="relu", name="dense_2")(x)
outputs = layers.Dense(10, activation="softmax", name="predictions")(x)

model = keras.Model(inputs=inputs, outputs=outputs)
```

```
In [35]: (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

# Preprocess the data (these are NumPy arrays)
x_train = x_train.reshape(60000, 784).astype("float32") / 255
x_test = x_test.reshape(10000, 784).astype("float32") / 255

y_train = y_train.astype("float32")
y_test = y_test.astype("float32")

# Reserve 10,000 samples for validation
x_val = x_train[-10000:]
y_val = y_train[-10000:]
x_train = x_train[:-10000]
y_train = y_train[:-10000]
```

```
In [36]: model.compile(
    optimizer=keras.optimizers.RMSprop(), # Optimizer
    # Loss function to minimize
    loss=keras.losses.SparseCategoricalCrossentropy(),
    # List of metrics to monitor
    metrics=[keras.metrics.SparseCategoricalAccuracy()],
)
```

```
In [37]: print("Fit model on training data")
history = model.fit(
    x_train,
    y_train,
    batch_size=64,
    epochs=2,
    # We pass some validation for
    # monitoring validation Loss and metrics
    # at the end of each epoch
    validation_data=(x_val, y_val),
)
```

```
Fit model on training data
Epoch 1/2
782/782 [=====] - 3s 3ms/step - loss: 0.3420 - sparse_categorical_accuracy: 0.9025 - val_loss: 0.1939 - val_sparse_categorical_accuracy: 0.9446
Epoch 2/2
782/782 [=====] - 1s 2ms/step - loss: 0.1616 - sparse_categorical_accuracy: 0.9521 - val_loss: 0.1493 - val_sparse_categorical_accuracy: 0.9545
```

```
In [38]: history.history
```

```
Out[38]: {'loss': [0.3419865667819977, 0.1616363227367401],
'sparse_categorical_accuracy': [0.9025200009346008, 0.9521399736404419],
'val_loss': [0.1939145028591156, 0.14933522045612335],
'val_sparse_categorical_accuracy': [0.944599986076355, 0.9545000195503235]}
```

```
In [39]: # Evaluate the model on the test data using `evaluate`
print("Evaluate on test data")
results = model.evaluate(x_test, y_test, batch_size=128)
print("test loss, test acc:", results)

# Generate predictions (probabilities -- the output of the last layer)
# on new data using `predict`
print("Generate predictions for 3 samples")
predictions = model.predict(x_test[:3])
print("predictions shape:", predictions.shape)
```

```
Evaluate on test data
79/79 [=====] - 0s 1ms/step - loss: 0.1532 - sparse_categorical_accuracy: 0.9512
test loss, test acc: [0.15317949652671814, 0.951200008392334]
Generate predictions for 3 samples
predictions shape: (3, 10)
```

```
In [40]: model.compile(
    optimizer=keras.optimizers.RMSprop(learning_rate=1e-3),
    loss=keras.losses.SparseCategoricalCrossentropy(),
    metrics=[keras.metrics.SparseCategoricalAccuracy()],
)
```

```
In [41]: model.compile(
    optimizer="rmsprop",
    loss="sparse_categorical_crossentropy",
    metrics=["sparse_categorical_accuracy"],
)
```

```
In [ ]:
```

```
In [42]: from sklearn import linear_model

# Create an instance of the classifier
classifier = linear_model.LogisticRegression()
```

```
In [43]: np.random.set_state(checkpoint_state)
options = {'c1': 0.5, 'c2': 0.5, 'w':0.9, 'k': SWARM_SIZE, 'p':2}
```

## final model

```
In [45]: # Initialize swarm, arbitrary
options = {'c1': 0.5, 'c2': 0.5, 'w':0.9, 'k': SWARM_SIZE, 'p':2}

# Call instance of PSO
dimensions = 15 # dimensions should be the number of features
optimizer = ps.discrete.BinaryPSO(n_particles=30, dimensions=dimensions, options=options)

# Perform optimization
cost, pos = optimizer.optimize(f, iters=1000, verbose=2)
```

```
2021-09-08 19:28:49,964 - pyswarms.discrete.binary - INFO - Optimize for 1000 iterations with {'c1': 0.5, 'c2': 0.5, 'w': 0.9, 'k': 30, 'p': 2}
pyswarms.discrete.binary: 100%|██████████| 1000/1000, best_cost=-36.6
2021-09-08 19:31:46,865 - pyswarms.discrete.binary - INFO - Optimization finished | best cost: -36.58769230769231, best pos: [0 0 0 0 0 1 0 0 0 0 0 0 0 1 0]
```

```
In [54]: # Create two instances of LogisticRegression
classifier = linear_model.LogisticRegression()

# Get the selected features from the final positions
X_selected_features = X[:,pos==1] # subset

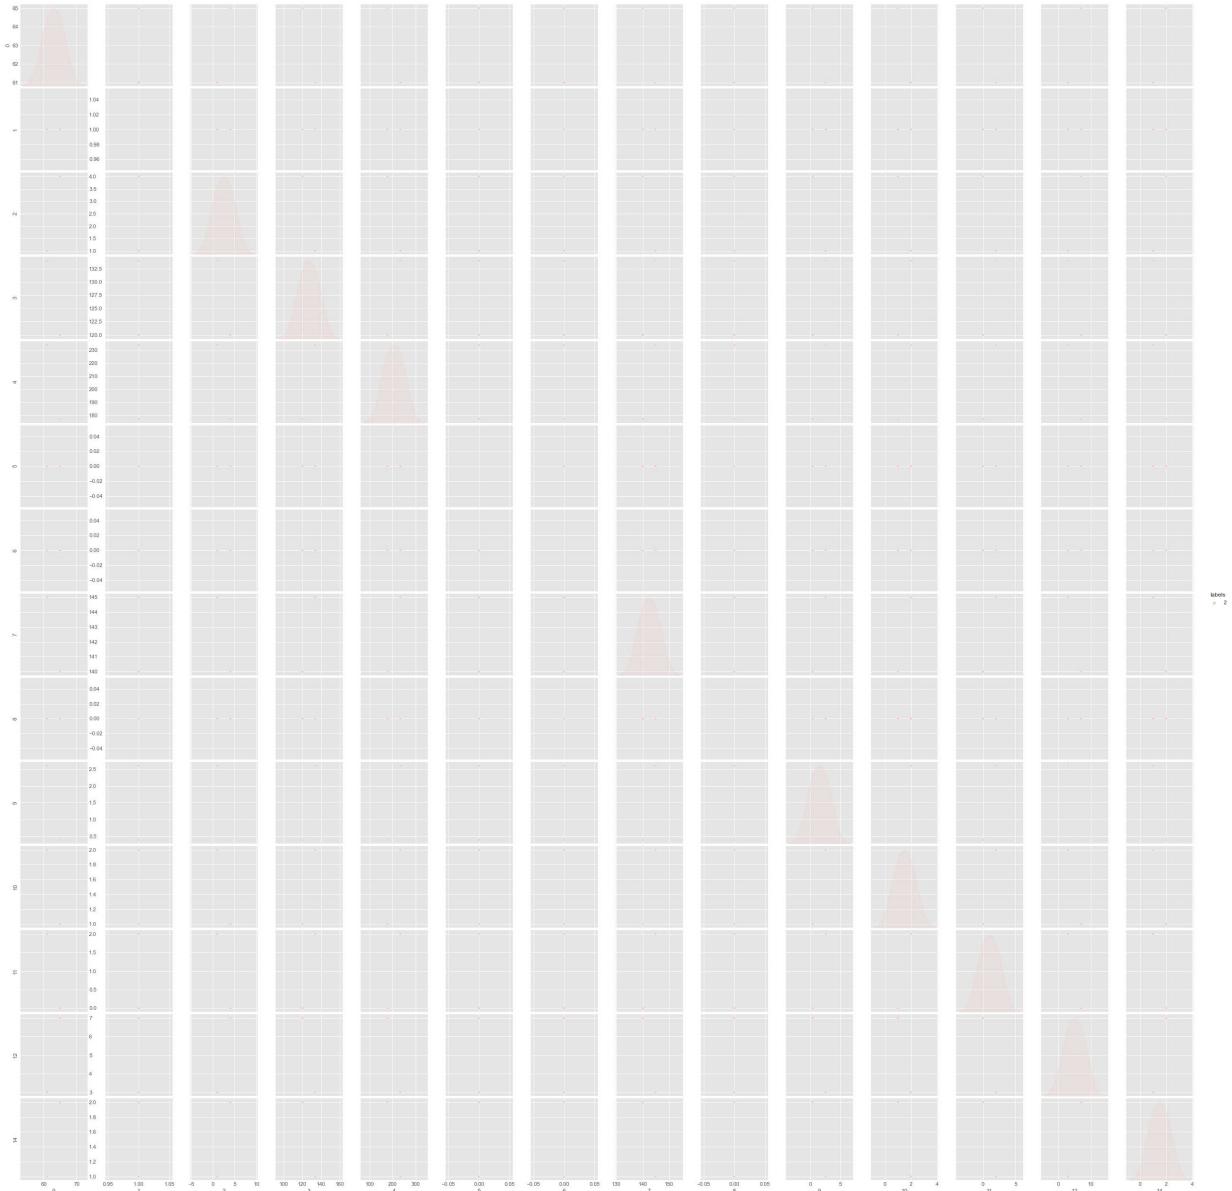
# Perform classification and store performance in P
classifier.fit(X_selected_features, y)
```

```
Out[54]: LogisticRegression()
```

```
In [72]: selected_features=df.values[np.where(pos==1)].tolist()
```

```
In [73]: df1 = pd.DataFrame(selected_features)
df1['labels'] = pd.Series(y)
sns.pairplot(df1,hue='labels');
```

```
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\distributions.py:306: UserWarning: Dataset has 0 variance; skipping density estimate.
    warnings.warn(msg, UserWarning)
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\distributions.py:306: UserWarning: Dataset has 0 variance; skipping density estimate.
    warnings.warn(msg, UserWarning)
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\distributions.py:306: UserWarning: Dataset has 0 variance; skipping density estimate.
    warnings.warn(msg, UserWarning)
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\distributions.py:306: UserWarning: Dataset has 0 variance; skipping density estimate.
    warnings.warn(msg, UserWarning)
```



```
In [119]: X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2)
```

```
In [120]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(X_train)

X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [121]: max(df['age'])
```

```
Out[121]: 77
```

```
In [ ]:
```

```
In [ ]:
```

```
In [75]: X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2)
```

```
In [81]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(X_train)

X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [87]: RandomForestClassifier, VotingClassifier, AdaBoostClassifier, GradientBoostingClassifier
```

## DECISION TREE

```
In [193]: from sklearn.tree import DecisionTreeClassifier
decc = DecisionTreeClassifier()
decc.fit(X_train,y_train)
y_pred_decc = decc.predict(X_test)
```

```
In [194]: from random import random
import pandas as pd
from matplotlib import pyplot as plt
```

```
In [195]: from sklearn.metrics import confusion_matrix
```

```
In [196]: CM=confusion_matrix(y_test,y_pred_decc)
sns.heatmap(CM, annot=True)

TN = CM[0][0]
FN = CM[1][0]
TP = CM[1][1]
FP = CM[0][1]
specificity = TN/(TN+FP)
print(specificity)
```

0.4



```
In [197]: from sklearn.metrics import accuracy_score
```

```
In [198]: acc= accuracy_score(y_test,y_pred_decc)
print(acc)
```

0.6

## BAGGED TREE

```
In [168]: from sklearn.ensemble import BaggingClassifier
```

```
In [169]: bt = BaggingClassifier(base_estimator=decc,n_estimators=270)
```

```
In [170]: bt.fit(X_train,y_train)
```

```
Out[170]: BaggingClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=270)
```

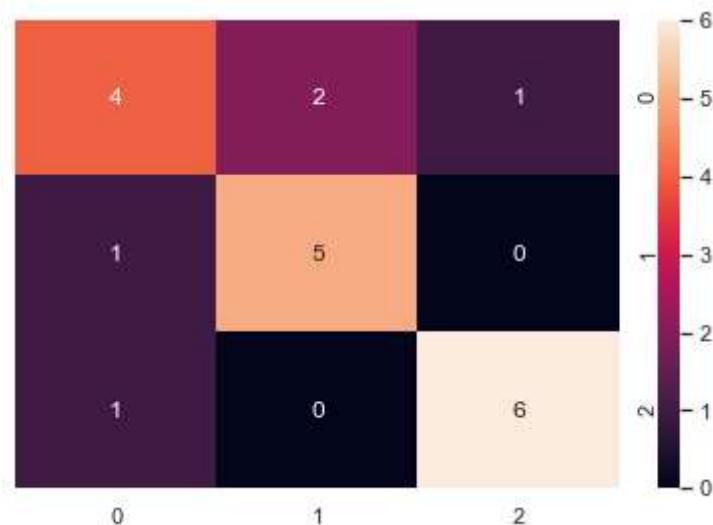
```
In [171]: y_pred_ada = bt.predict(X_test)
```

```
In [172]: from sklearn.metrics import log_loss,roc_auc_score,precision_score,f1_score,recall_score,classification_report,confusion_matrix,accuracy_score
from sklearn import metrics
```

```
In [173]: CM=confusion_matrix(y_test,y_pred_ada)
sns.heatmap(CM, annot=True)
```

```
TN = CM[0][0]
FN = CM[1][0]
TP = CM[1][1]
FP = CM[0][1]
specificity = TN/(TN+FP)
print(specificity)
```

0.6666666666666666



```
In [174]: acc= accuracy_score(y_test,y_pred_ada)
print(acc)
```

0.75

## RANDOM FOREST

```
In [230]: rf_ent = RandomForestClassifier(criterion='entropy',n_estimators=100)
rf_ent.fit(X_train, y_train)
y_pred_rfe = rf_ent.predict(X_test)
```

```
In [231]: CM=confusion_matrix(y_test,y_pred_rfe)
sns.heatmap(CM, annot=True)

TN = CM[0][0]
FN = CM[1][0]
TP = CM[1][1]
FP = CM[0][1]
specificity = TN/(TN+FP)
print(specificity)
```

0.6



```
In [232]: acc= accuracy_score(y_test,y_pred_rfe)
print(acc)
```

0.65

## KNN

```
In [243]: from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(X, y)
```

```
Out[243]: KNeighborsClassifier()
```

```
In [244]: y_pred = classifier.predict(X_test)
```

```
In [245]: from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_test,y_pred))
print(classification_report(y_test,y_pred))
```

```
[[4 1 2]
 [1 5 0]
 [0 1 6]]
          precision    recall  f1-score   support

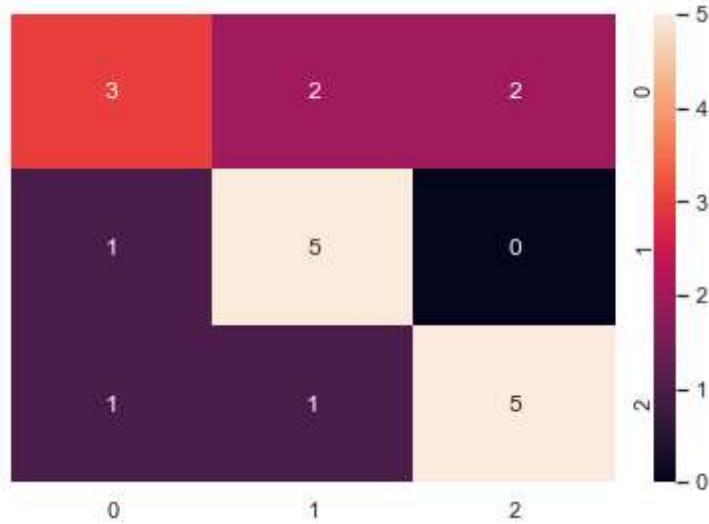
           0       0.80      0.57      0.67       7
           1       0.71      0.83      0.77       6
           2       0.75      0.86      0.80       7

    accuracy                           0.75      20
   macro avg       0.75      0.75      0.75      20
weighted avg       0.76      0.75      0.74      20
```

```
In [246]: CM=confusion_matrix(y_test,y_pred_rfe)
sns.heatmap(CM, annot=True)
```

```
TN = CM[0][0]
FN = CM[1][0]
TP = CM[1][1]
FP = CM[0][1]
specificity = TN/(TN+FP)
print(specificity)
```

0.6



```
In [247]: acc= accuracy_score(y_test,y_pred)
print(acc)
```

0.75

## ADABOOST

```
In [254]: adaboost = AdaBoostClassifier()  
adaboost.fit(X_train,y_train)  
y_pred_adaboost = adaboost.predict(X_test)
```

```
In [255]: acc= accuracy_score(y_test,y_pred_adaboost)  
print(acc)
```

0.55

```
In [ ]:
```

```
In [ ]:
```

In [20]: `df.value_counts()[:20]`

Out[20]:

age	sex	chest	resting_blood_pressure	serum_cholesterol	fasting_blood_sug	ar	resting_electrocardiographic_results	maximum_heart_rate_achieved	exerci	se_induced_angina	oldpeak	slope	number_of_major_vessels	thal	class
29	1	2	130		204					0					0
2					202										0
0.0		1	0		3	absent				1					0
59	1	4	170		326										1
2					140										1
3.4		3	0		7	present				1					0
58	1	4	128		259										0
2					130										1
3.0		2	2		7	present				1					0
			146		218										0
0					105										0
2.0		2	1		7	present				1					0
			150		270										0
2					111										1
0.8		1	0		7	present				1					0
59	0	4	174		249										1
0					143										1
0.0		2	0		3	present				1					0
	1	1	160		273										0
2					125										0
0.0		1	0		3	present				1					0
			170		288										0
2					159										0
0.2		2	0		7	present				1					0
			178		270										0
2					145										0
4.2		3	0		7	absent				1					0
		2	140		221										0
0					164										1
0.0		1	0		3	absent				1					1
		3	126		218										1
0					134										0
2.2		2	1		6	present				1					0
			150		212										1
0					157										0
1.6		1	0		3	absent				1					0
		4	110		239										0
2					142										1
1.2		2	1		7	present				1					0
			135		234										0
0					161										0
0.5		2	0		7	absent				1					0
			138		271										0
2					182										0
0.0		1	0		3	absent				1					0
			140		177										0
0					162										1
0.0		1	1		7	present				1					0
60	0	1	150		240										0
0					171										0
0.9		1	0		3	absent				1					0

```

58    1    4    125          300        0
2
0.0    1    2          171        0
60    0    3   102          7    present    1
0           318        0
0           160        0
0.0    1    1          3    absent     1
0           120          178        1
0           96        0
0.0    1    0          3    absent     1
dtype: int64

```

In [22]:

```
X = df.drop(['class'], axis=1)
y = df['class']
```

In [122]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2)
```

In [123]:

```
print('Distribution of target variable in training set')
print(y_train.value_counts())

print('Distribution of target variable in test set')
print(y_test.value_counts())
```

Distribution of target variable in training set

---

**AttributeError** Traceback (most recent call last)  
<ipython-input-123-007fda9da671> in <module>  
 1 print('Distribution of target variable in training set')  
----> 2 print(y\_train.value\_counts())  
 3  
 4 print('Distribution of target variable in test set')  
 5 print(y\_test.value\_counts())

**AttributeError**: 'numpy.ndarray' object has no attribute 'value\_counts'

In [124]:

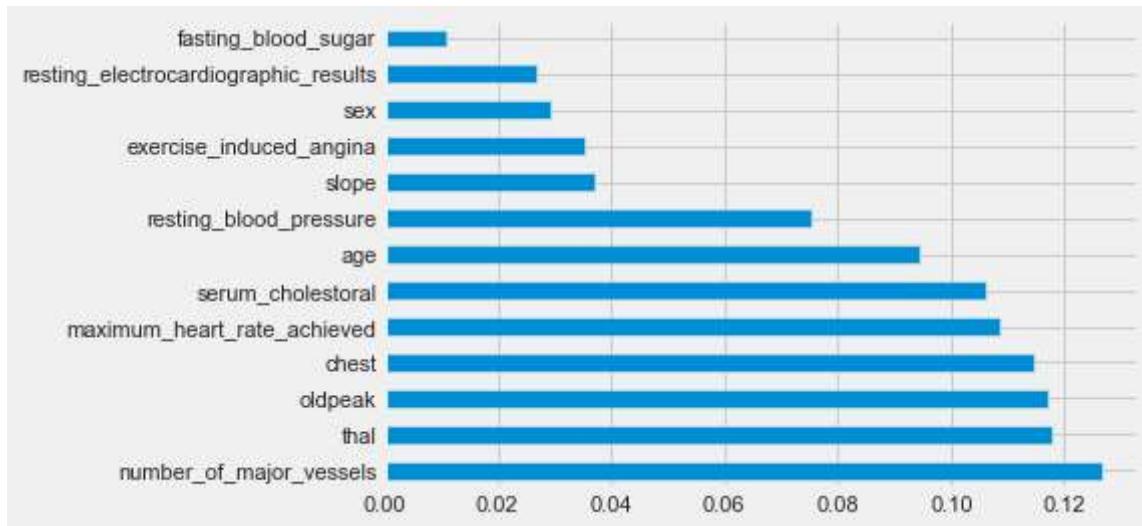
```
print('-----Training Set-----')
print(X_train.shape)
print(y_train.shape)

print('-----Test Set-----')
print(X_test.shape)
print(y_test.shape)
```

-----Training Set-----  
(80, 15)  
(80,)  
-----Test Set-----  
(20, 15)  
(20,)

```
In [43]: feat_importances = pd.Series(rf_ent.feature_importances_, index=X_train.columns)
feat_importances.nlargest(20).plot(kind='barh')
```

Out[43]: <AxesSubplot:>



```
In [49]: X = df.iloc[:,0:12]
y = df.iloc[:,12]
```

```
In [51]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st
```

```
In [52]: rt = DecisionTreeRegressor(criterion = 'mse', max_depth=5)
```

```
In [53]: rt.fit(X_train,y_train)
```

Out[53]: DecisionTreeRegressor(max\_depth=5)

```
In [54]: y_pred = rt.predict(X_test)
r2_score(y_test,y_pred)
```

Out[54]: 0.10740364193337304

```
In [55]: param_grid = {
    'max_depth':[2,4,8,10,None],
    'criterion':['mse','mae'],
    'max_features':[0.25,0.5,1.0],
    'min_samples_split':[0.25,0.5,1.0]
}
```

```
In [56]: reg = GridSearchCV(DecisionTreeRegressor(),param_grid=param_grid)
```

```
In [57]: reg.fit(X_train,y_train)
```

```
Out[57]: GridSearchCV(estimator=DecisionTreeRegressor(),
                       param_grid={'criterion': ['mse', 'mae'],
                                   'max_depth': [2, 4, 8, 10, None],
                                   'max_features': [0.25, 0.5, 1.0],
                                   'min_samples_split': [0.25, 0.5, 1.0]})
```

```
In [36]: reg.best_score_
```

```
Out[36]: 0.16345729494299172
```

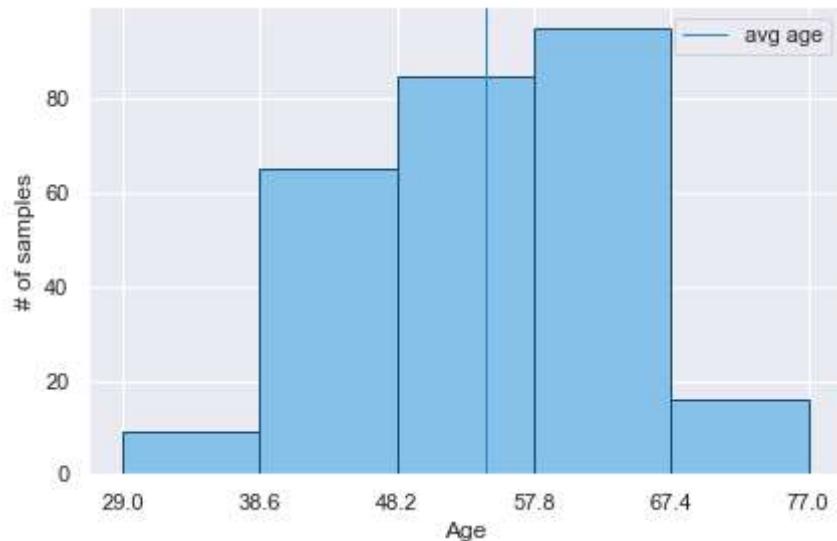
```
In [38]: reg.best_params_
```

```
Out[38]: {'criterion': 'mse',
           'max_depth': 2,
           'max_features': 0.25,
           'min_samples_split': 1.0}
```

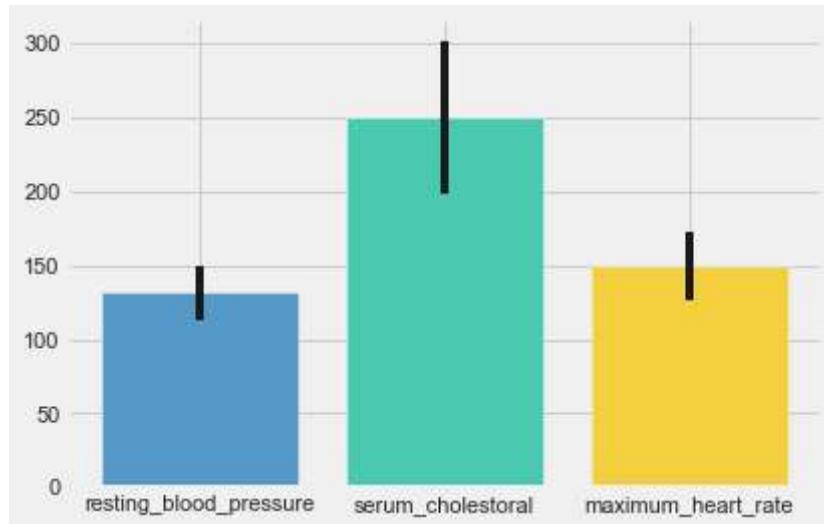
```
In [25]: count, bin_edges = np.histogram(df['age'], bins = 5)
```

```
In [ ]:
```

```
In [26]: age_mean = df['age'].mean()
plt.hist(df['age'], color = '#85C1E9', ec = '#1B4F72', bins = 5)
plt.xlabel('Age')
plt.ylabel('# of samples')
plt.xticks(bin_edges)
plt.axvline(age_mean, label = 'avg age',
            color = '#2E86C1', linewidth = 1)
plt.legend()
plt.show()
```



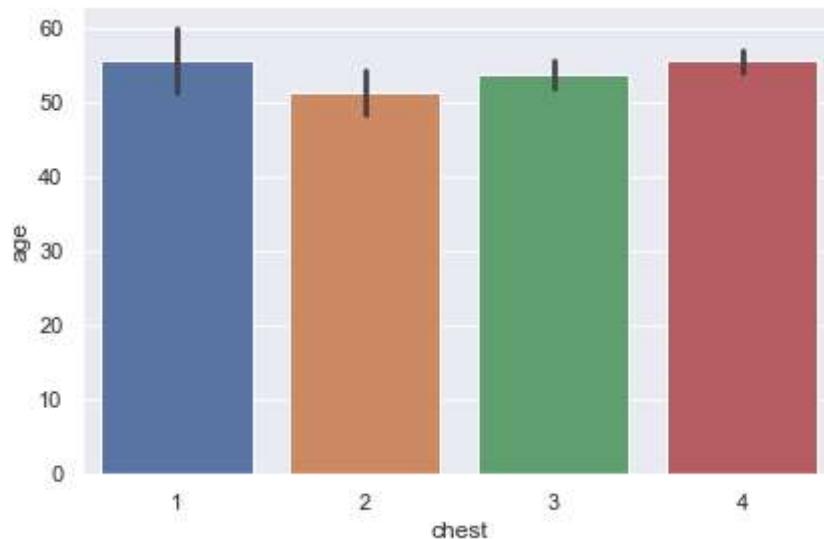
```
In [20]: plt.bar('resting_blood_pressure', df['resting_blood_pressure'].mean(), yerr = df['resting_blood_pressure'].std(), color = colors[1])
plt.bar('serum_cholestorol', df['serum_cholestorol'].mean(), yerr = df['serum_cholestorol'].std(), color = colors[3])
plt.bar('maximum_heart_rate', df['maximum_heart_rate_achieved'].mean(), yerr = df['maximum_heart_rate_achieved'].std(), color = colors[5])
plt.show()
```



```
In [27]: sns.barplot(df['chest'], df['age'])
```

C:\Users\Priya\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

```
Out[27]: <AxesSubplot:xlabel='chest', ylabel='age'>
```

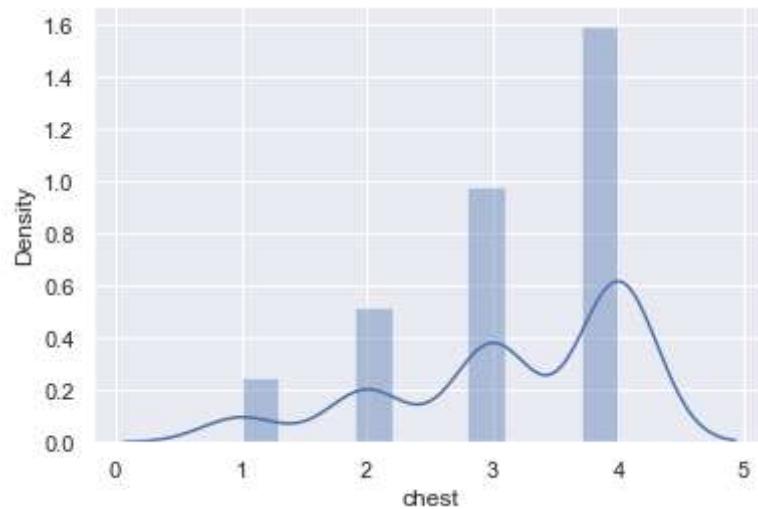


In [36]: `sns.distplot(df['chest'])`

```
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
```

```
warnings.warn(msg, FutureWarning)
```

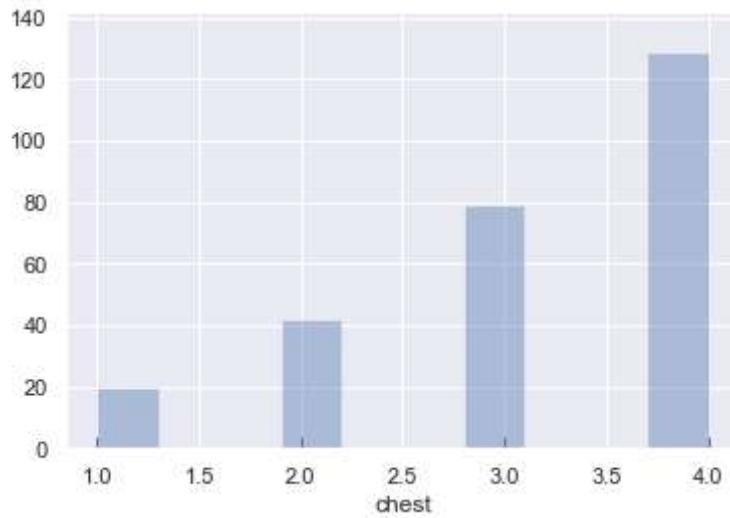
Out[36]: <AxesSubplot:xlabel='chest', ylabel='Density'>



In [37]: `sns.distplot(df['chest'], kde=False, rug=True)`

```
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\distributions.py:2056: FutureWarning: The `axis` variable is no longer used and will be removed. Instead, assign variables directly to `x` or `y`.
  warnings.warn(msg, FutureWarning)
```

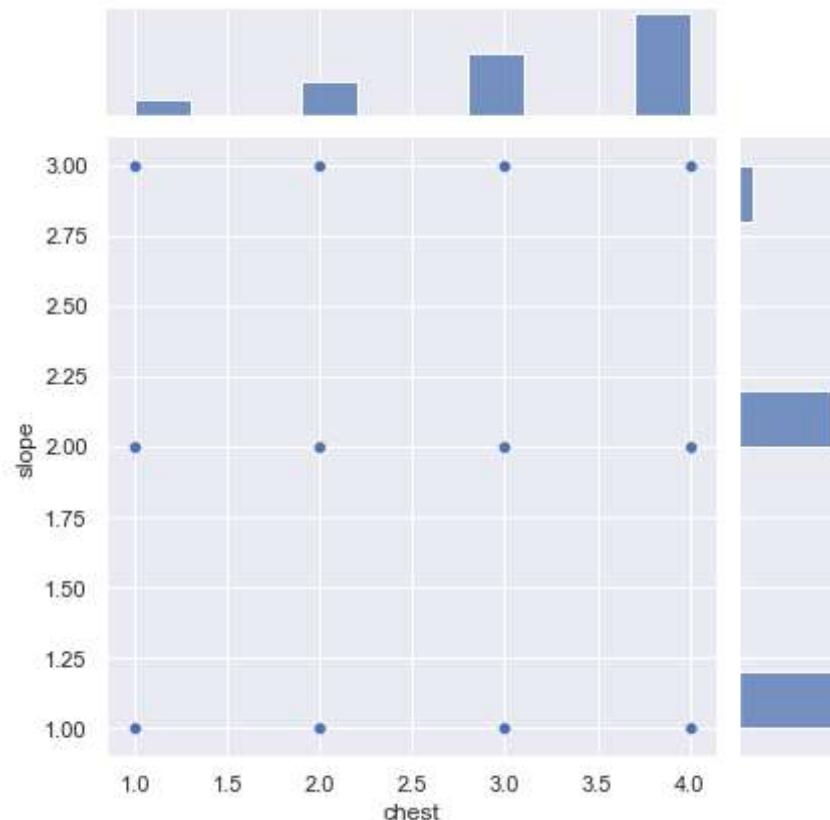
Out[37]: <AxesSubplot:xlabel='chest'>



In [38]: `sns.jointplot(df['chest'], df['slope'])`

C:\Users\Priya\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

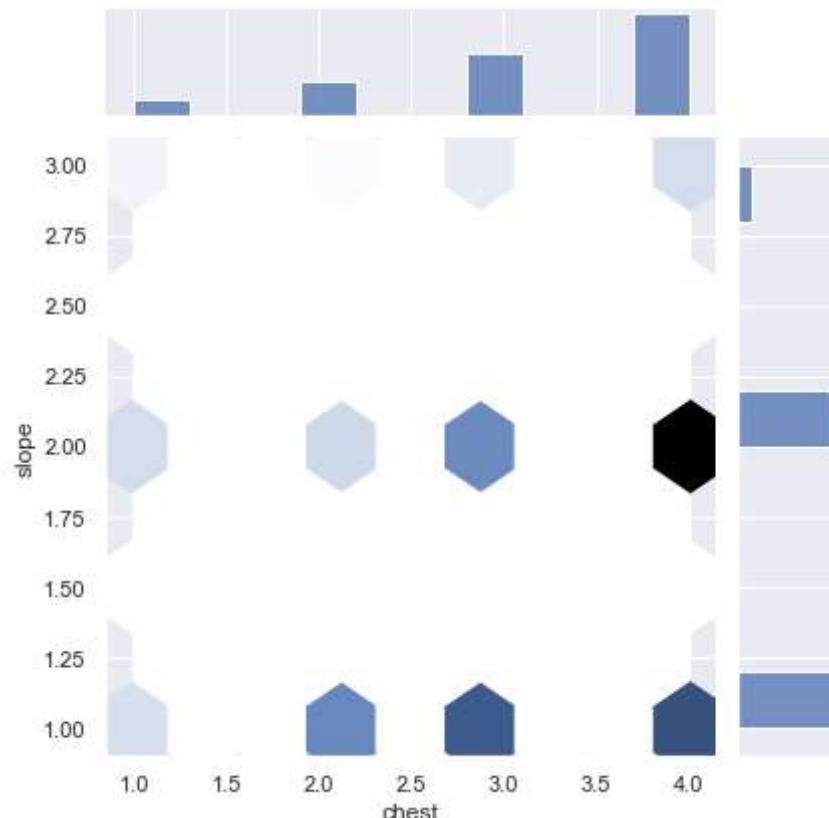
Out[38]: <seaborn.axisgrid.JointGrid at 0x1d86d103160>



```
In [39]: sns.jointplot(df['chest'], df['slope'], kind="hex")
```

C:\Users\Priya\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

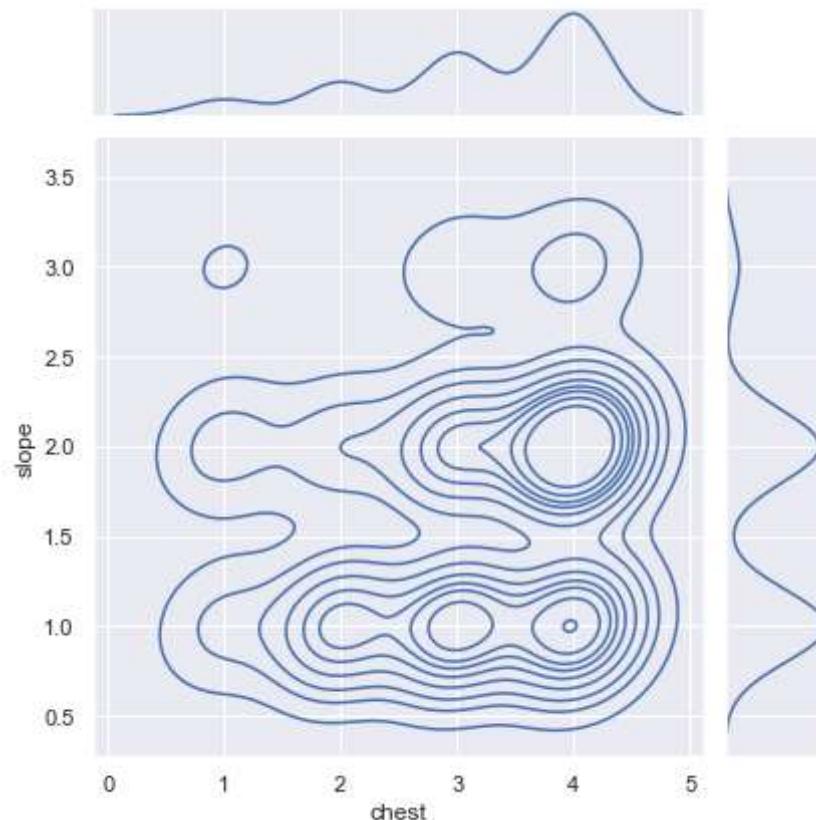
```
Out[39]: <seaborn.axisgrid.JointGrid at 0x1d86d0223d0>
```



```
In [40]: sns.jointplot(df['chest'], df['slope'], kind="kde")
```

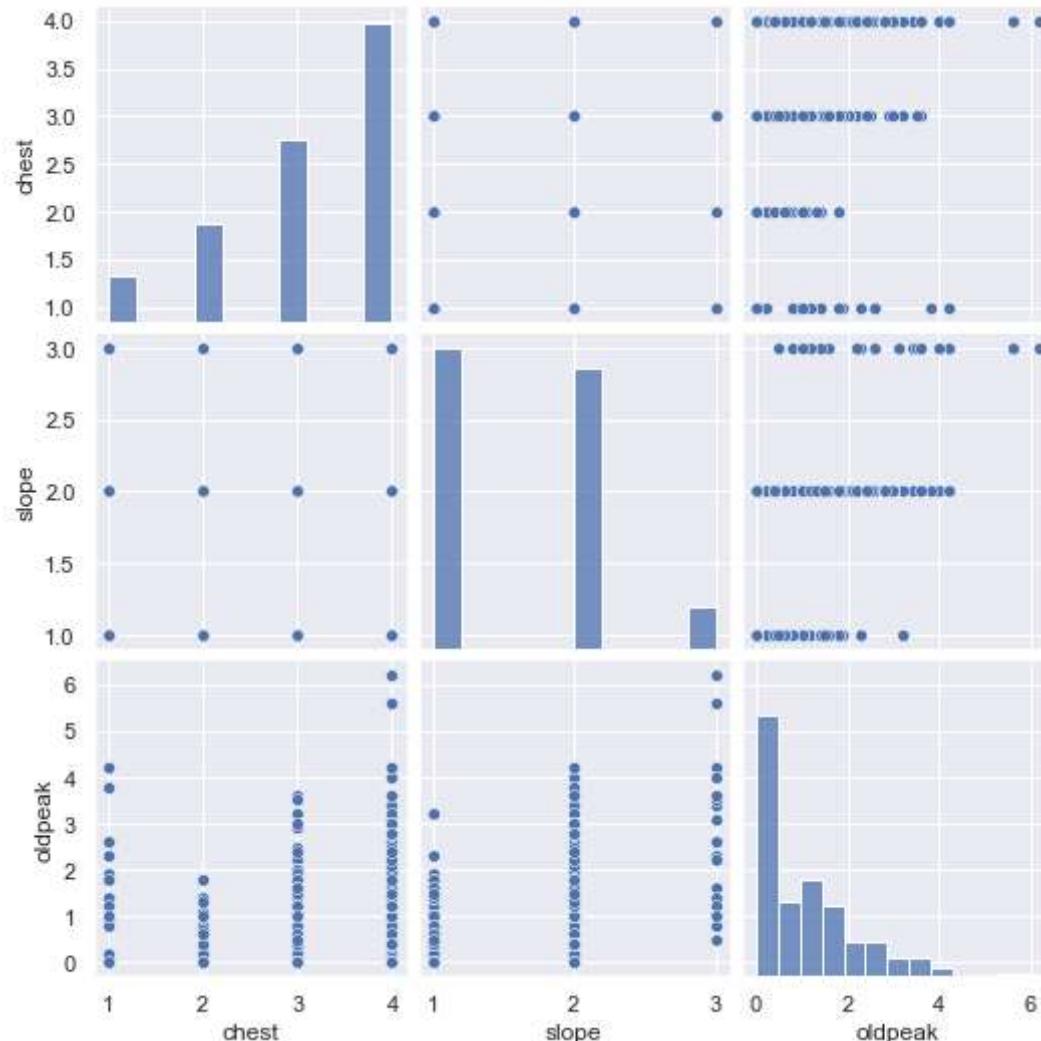
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

```
Out[40]: <seaborn.axisgrid.JointGrid at 0x1d86d14aa90>
```



```
In [41]: sns.pairplot(df[['chest', 'slope', 'oldpeak']])
```

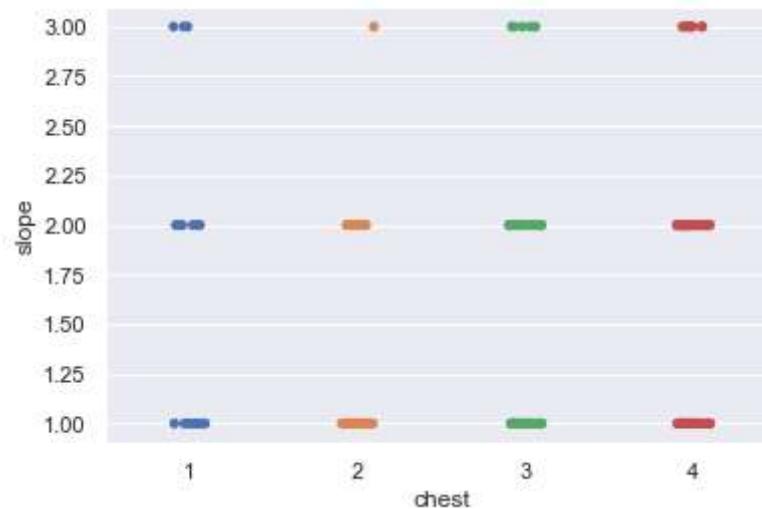
```
Out[41]: <seaborn.axisgrid.PairGrid at 0x1d86e3d14f0>
```



```
In [42]: sns.stripplot(df['chest'], df['slope'])
```

C:\Users\Priya\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

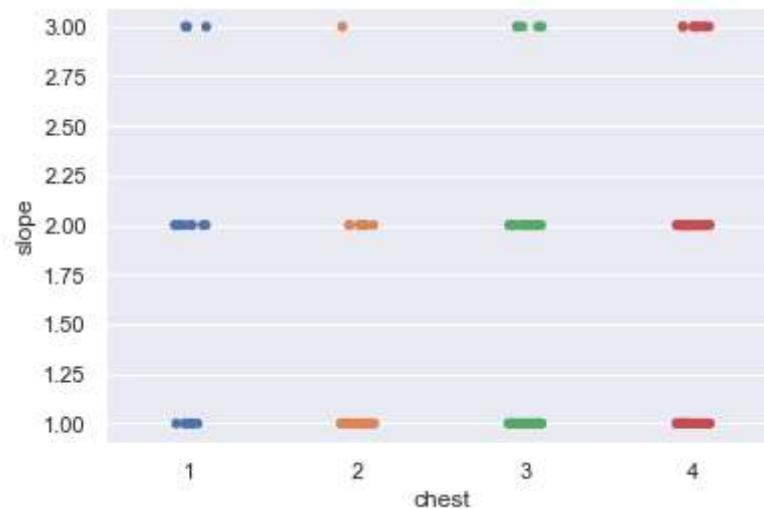
```
Out[42]: <AxesSubplot:xlabel='chest', ylabel='slope'>
```



```
In [43]: sns.stripplot(df['chest'], df['slope'], jitter = True)
```

C:\Users\Priya\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

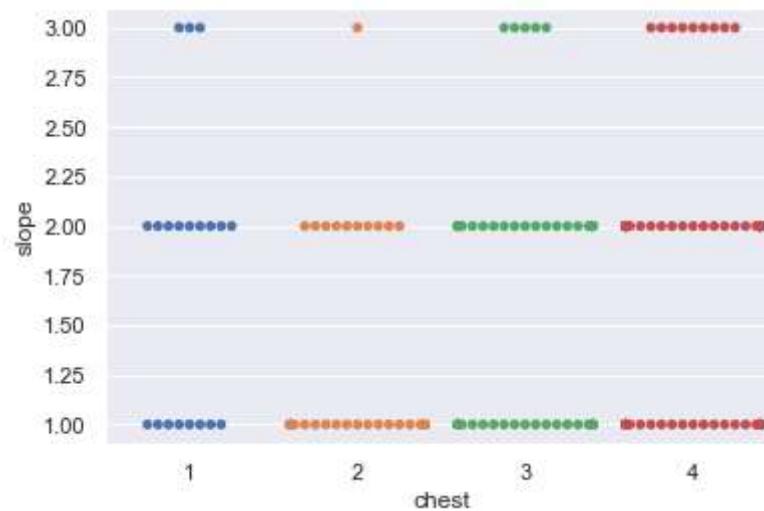
```
Out[43]: <AxesSubplot:xlabel='chest', ylabel='slope'>
```



In [44]: `sns.swarmplot(df['chest'], df['slope'])`

```
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
    warnings.warn(  
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning:  
    42.9% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.  
    warnings.warn(msg, UserWarning)  
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning:  
    60.8% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.  
    warnings.warn(msg, UserWarning)  
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning:  
    72.9% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.  
    warnings.warn(msg, UserWarning)
```

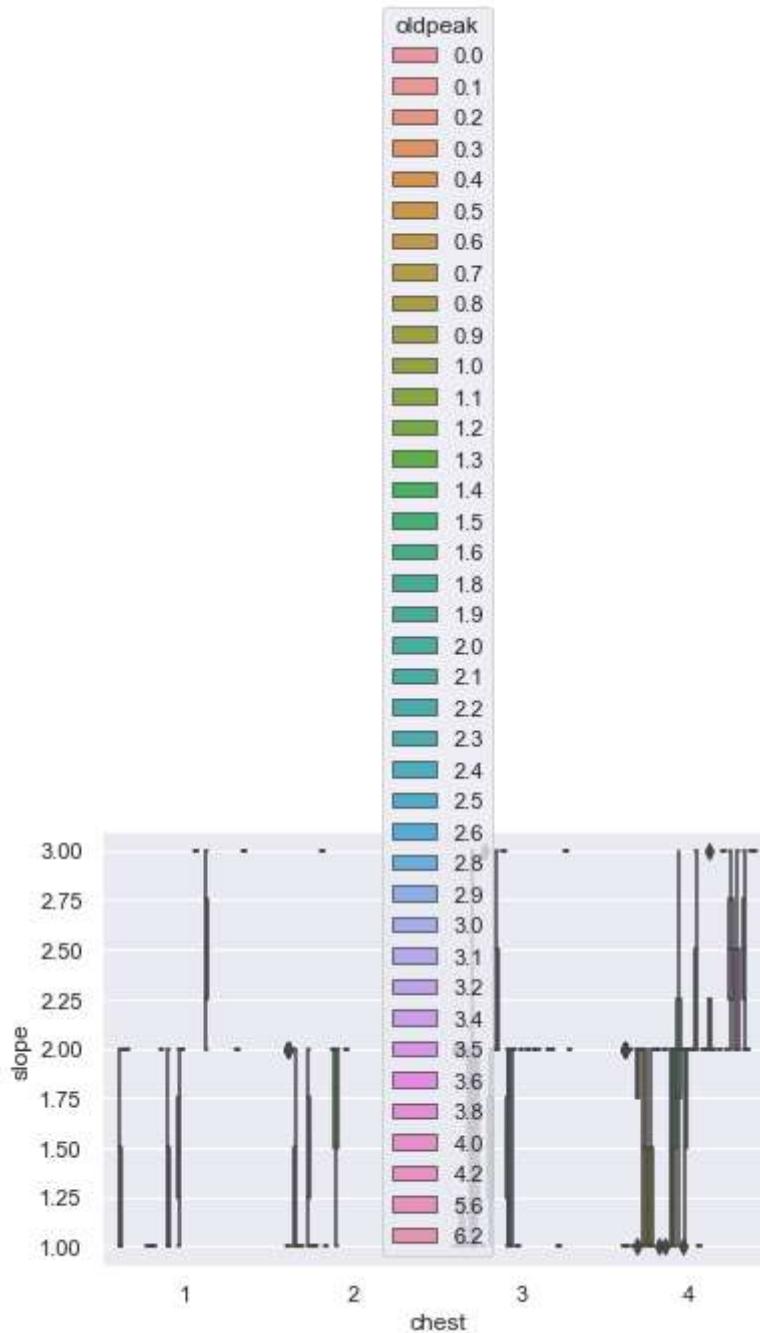
Out[44]: <AxesSubplot:xlabel='chest', ylabel='slope'>



```
In [45]: sns.boxplot(df['chest'], df['slope'], hue=df['oldpeak'])
```

C:\Users\Priya\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

```
Out[45]: <AxesSubplot:xlabel='chest', ylabel='slope'>
```

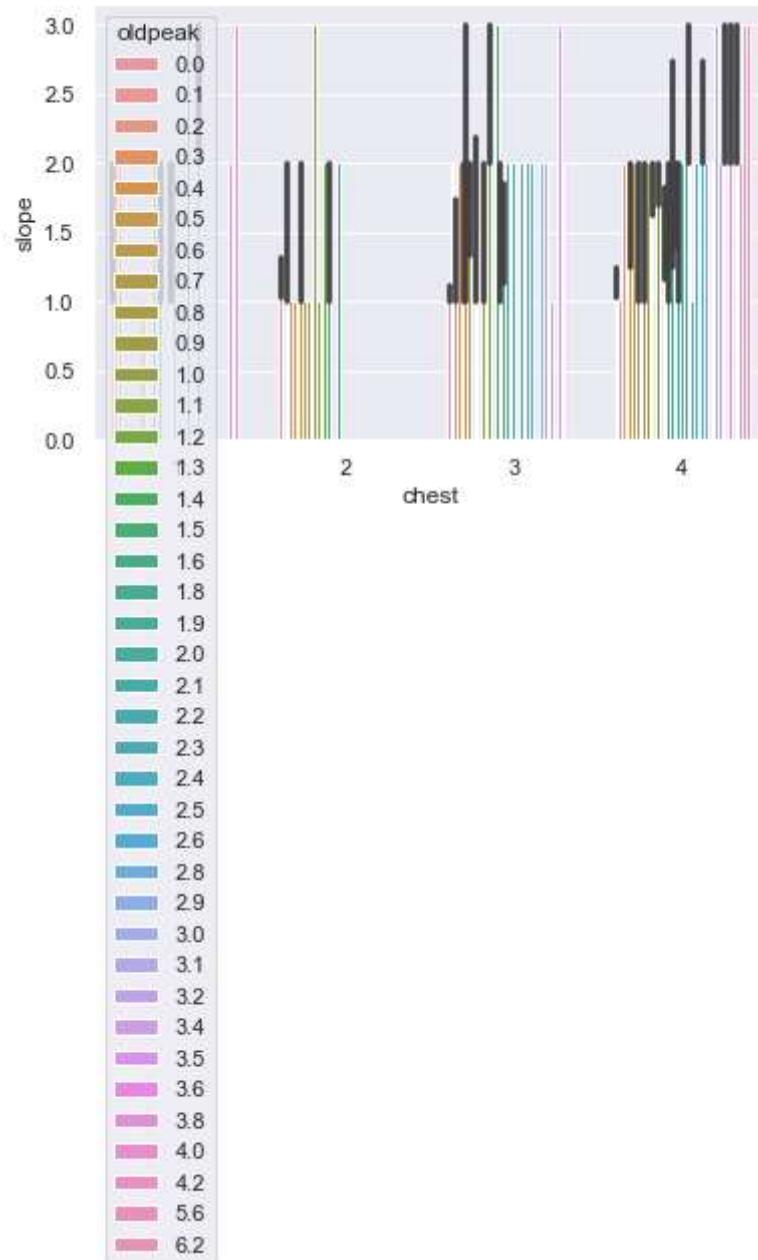




```
In [46]: sns.barplot(df['chest'], df['slope'], hue=df['oldpeak'])
```

C:\Users\Priya\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

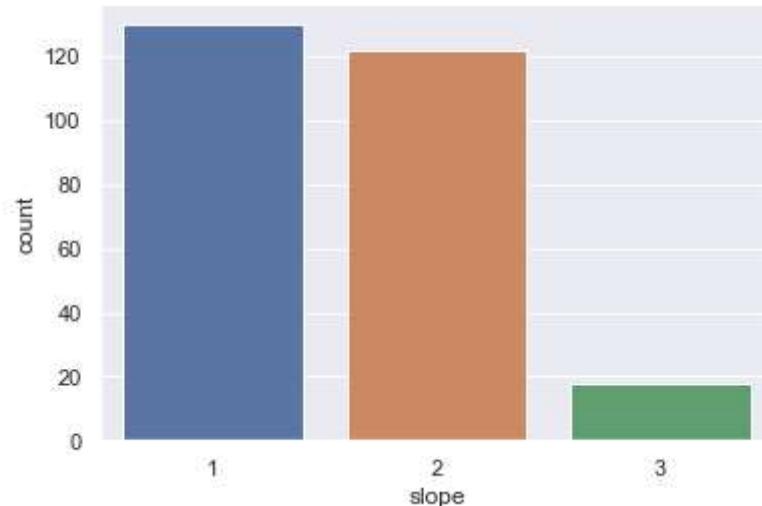
```
Out[46]: <AxesSubplot:xlabel='chest', ylabel='slope'>
```



In [47]: `sns.countplot(df['slope'])`

```
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(
```

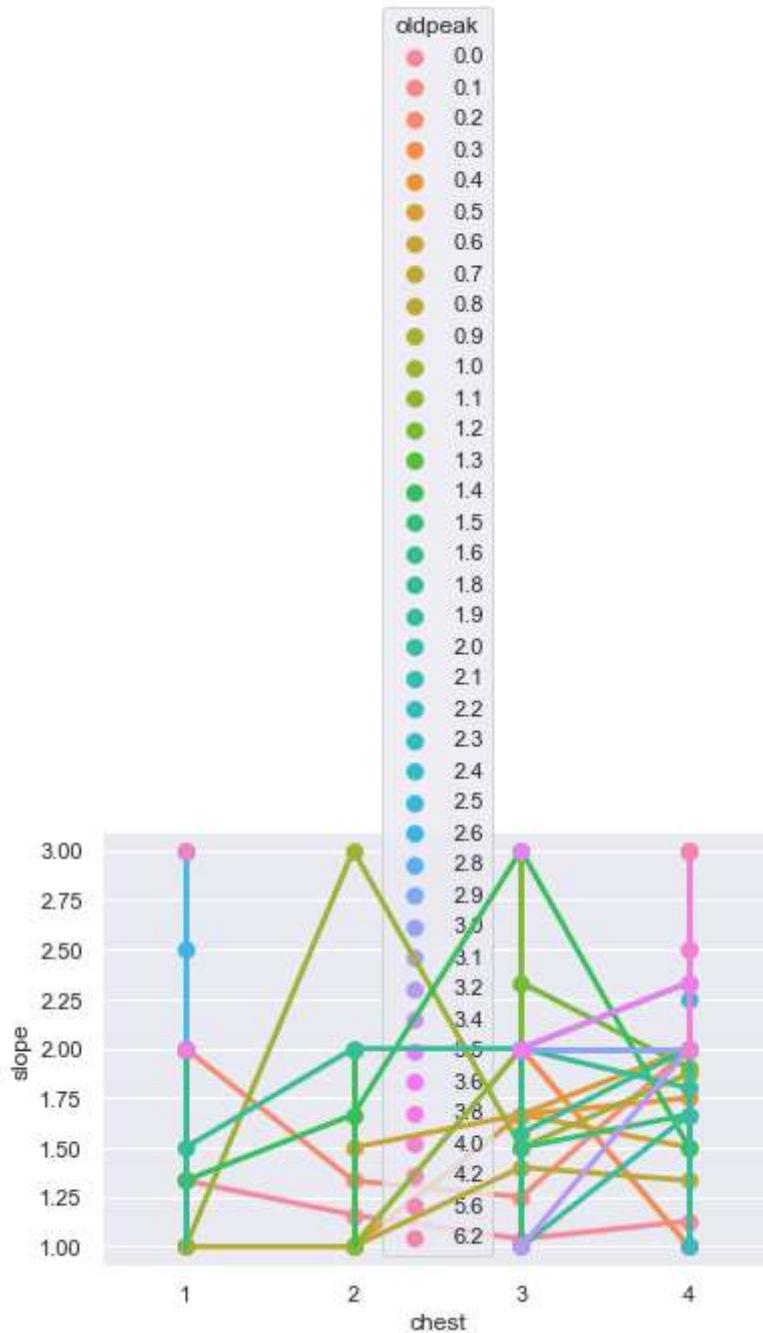
Out[47]: <AxesSubplot:xlabel='slope', ylabel='count'>



```
In [48]: sns.pointplot(df['chest'], df['slope'], hue=df['oldpeak'])
```

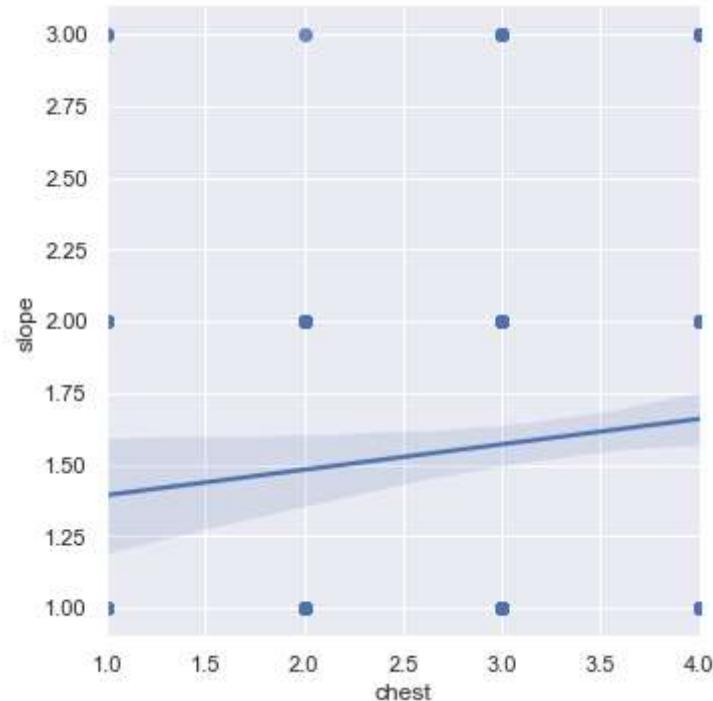
C:\Users\Priya\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

```
Out[48]: <AxesSubplot:xlabel='chest', ylabel='slope'>
```



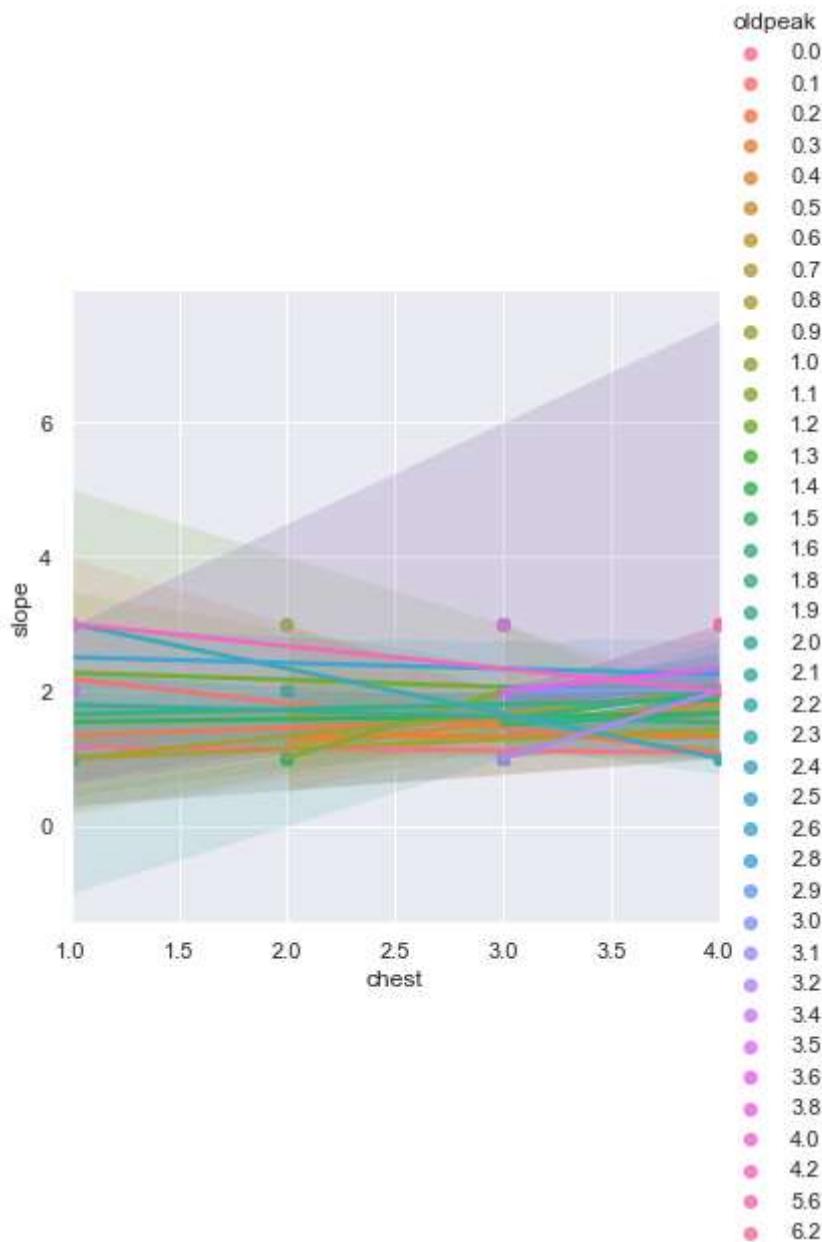
```
In [49]: sns.lmplot(x="chest", y="slope", data=df)
```

```
Out[49]: <seaborn.axisgrid.FacetGrid at 0x1d86faca5b0>
```



```
In [50]: sns.lmplot(x="chest", y="slope", hue="oldpeak", data=df)
```

```
Out[50]: <seaborn.axisgrid.FacetGrid at 0x1d86f783e50>
```



In [ ]:

In [ ]: